Yang Sijun

AI Engineer | Male | 23

yangsijun233@gmail.com | +86 18207579044 | github.com/2499146834-arch

Education

Lingnan University, Hong Kong — M.Sc. in Data Science *2025 – 2026*

Core courses: Machine Learning, Deep Learning, Computer Vision, Data Mining, Artificial Intelligence

Guangzhou Maritime University — B.A. in Economics and Management *2021 – 2025*

Skills

- **Engineering:** Proficient in Python, PyTorch, TensorFlow; daily use of Claude Code for AI-assisted development; fine-tune and deploy 7B LoRA models on a single GPU via LLaMA-Factory; hands-on with ResNet / YOLO / MobileNetV2 transfer learning, mixed-precision training and GPU memory optimization.
- **Academic:** Major in AI / ML; familiar with knowledge distillation, chain-of-thought (CoT), SFT, attention mechanisms and collaborative filtering; experienced with LoRA, cosine annealing, progressive ablation and multi-seed reproducibility studies.
- **Languages:** Mandarin and Cantonese (native); English (IELTS 6.0).

Projects

FinBridge — Knowledge Distillation for Financial LLMs *2026.02 – 2026.05*

Distilled financial reasoning from a commercial LLM into a 7B model (Qwen2.5) via LoRA fine-tuning on LLaMA-Factory. Designed the student to emit executable "computation programs" rather than direct answers, sharply reducing hallucinations, and ran 6 controlled ablations to quantify each training stage. FinQA accuracy improved from 22% to **65.39%**, with auditable step-by-step outputs and single-GPU deployment on an RTX 5060 Ti.

Cats vs. Dogs Image Recognition — Two-Stage Transfer Learning *2026.02 – 2026.03*

Original CNN trained from scratch reached only 89% with errors concentrated in low-light scenes; led the rearchitecture for production. Switched to MobileNetV2 with a "freeze-backbone then fine-tune" two-stage strategy and added targeted low-light augmentation. Accuracy rose to **98.98%**, model size shrank ~9x, far better suited for mobile deployment.

Real-Time Parking Detection — Dual-YOLO Fine-Tuning *2026.02 – 2026.03*

Smart-parking scenes mix fixed CCTV with handheld phone shots, which no single model handled well. Designed a "specialist + generalist" dual-YOLO setup, switched to true-scale fine-tuning in the final stage, and fixed recurring VRAM crashes during training. Both models exceeded 99.4% accuracy and 99.7% precision/recall — **a 5.4-point gain over the original paper**, meeting industrial real-time requirements.

AODL — CIFAR-100 Image Classification (ResNet-18 + SE Attention) *2026.01 – 2026.02*

Original Top-1 accuracy was 72% with architectural choices unvalidated by controls. Reproduced and improved the pipeline in PyTorch with cosine annealing, label smoothing and mixed-precision training stacked, and ran 3-seed controls on the SE module. Top-1 rose to **75.66%** with 1.5–2x faster training; controls showed the SE block gave no significant gain on this task, correcting the original claim.

Internships

Qiye (Foshan) Software Co., Ltd. | Recommendation Algorithm Intern *2023.07 – 2023.09*

- **Responsibilities:** Tackled two client problems: a data center chiller tuned by operator intuition with poor energy efficiency, and an internal content platform suffering severe new-user churn and low CTR. Diagnosed both as "unverifiable black-box decisions"; applied association-rule mining for chiller operating conditions and upgraded collaborative filtering with a repaired offline evaluation pipeline.
- **Outcomes:** Identified the most stable chiller temperature band as the basis for energy-saving scheduling; for recommendations, pinpointed insufficient cold-start exposure as the root cause, fixed a metric-inflation flaw in offline evaluation, and benchmarked CF algorithms for a clear accuracy lift.

Foshan Yixin Network Technology Platform | Big Data Analytics Intern *2024.07–2024.09*

- **Responsibilities:** An e-commerce client could not process the daily flood of user reviews in time, so negative sentiment was only caught after escalation. Treating star ratings as a quantitative anchor for user sentiment, framed it as automatic "review text → star rating" classification and built the full data pipeline end to end.
- **Outcomes:** Delivered a reusable star-rating classification pipeline; identified and fixed a data leak inflating offline metrics, then benchmarked several mainstream models via cross-validation to select the best, which held up robustly on a strict test split.